

# CSC7072: Databases, fall 2015

Dr. Kim Bouters



nested queries in SQL, derived relations, and views

# nested queries, derived relations, and views

how to retrieve and manipulate data from a DB

---

we will be looking at:

- basic queries (e.g. selecting data, linking tables, sorting)
- set operations (e.g. joins, union, difference, intersection)
- aggregate functions (e.g. average, minimum, sum)
- null values (e.g. handling missing information)
- complex queries (*i.e.* putting it all together)
- modifying data
- nested subqueries (*i.e.* a query as part of a query)
- joins and views

# nested queries, derived relations, and views

retrieving data using SQL

---

don't forget:

```
SELECT {attribute [AS new_attribute_name]}  
      FROM {table [AS new_table_name]}  
          [{JOIN table ON attribute = attribute}]  
          [WHERE {condition}]  
          [GROUP BY {attribute}]  
          [HAVING {condition}]  
          [ORDER BY {attribute}]
```


where {argument} denotes you need to have at least one, and  
where [argument] denotes a part that is optional and can be omitted

# nested queries, derived relations, and views


retrieving data using SQL

---


remember these set operations?



courses thought in  
either 2014 or 2015



courses thought in  
both 2014 and 2015



courses only thought  
in 2014, not in 2015

all courses thought in 2014:

```
(SELECT course_id FROM section WHERE year = 2014)
```

all courses thought in 2015:


```
(SELECT course_id FROM section WHERE year = 2015)
```

# nested queries, derived relations, and views

## retrieving data using SQL

---


remember these set operations?



(SELECT course\_id FROM section WHERE year = 2014)

UNION


(SELECT course\_id FROM section WHERE year = 2015)



(SELECT course\_id FROM section WHERE year = 2014)

INTERSECT

(SELECT course\_id FROM section WHERE year = 2015)



(SELECT course\_id FROM section WHERE year = 2014)

EXCEPT

(SELECT course\_id FROM section WHERE year = 2015)

# nested queries, derived relations, and views

retrieving data using SQL

---

remember these set operations?

UNION, INTERSECT, EXCEPT will remove all duplicates

can be overridden by adding ALL at the end:

UNION ALL

INTERSECT ALL

EXCEPT ALL

how do you write a query to find all instructors working in both the computer science and the psychology department?

# nested queries, derived relations, and views

## retrieving data using SQL

---

many other set operations possible:

### problem

find the instructors earning more than some/all other instructors in the computer science department (CS due to space constraints)

```
SELECT name
FROM instructor
WHERE salary > SOME(
    SELECT salary
    FROM instructor
    WHERE dept_name = 'CS')
```

```
SELECT name
FROM instructor
WHERE salary > ALL(
    SELECT salary
    FROM instructor
    WHERE dept_name = 'CS')
```

# nested queries, derived relations, and views

## retrieving data using SQL

---

how does  $>\text{SOME}(\dots)$  work?

$5 > \text{SOME}(\begin{array}{|c|} \hline 2 \\ \hline 5 \\ \hline 6 \\ \hline \end{array})$       *true*, because  $5 > 2$

$5 > \text{SOME}(\begin{array}{|c|} \hline 9 \\ \hline 5 \\ \hline 6 \\ \hline \end{array})$       *false*, because  $5 \not> 5, 5 \not> 6, 5 \not> 9$

$5 = \text{SOME}(\begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array})$       *true*, because  $5 = 5$ , equal to "5 in(...)"

$5 <> \text{SOME}(\begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array})$       *true*, because  $5 <> 0$ , **not** equal to "5 not in(...)"



# nested queries, derived relations, and views

retrieving data using SQL

---

still many other set operations possible:

```
SELECT name
```

```
FROM instructor
```

```
WHERE salary IN(
```

```
    SELECT salary
```

```
    FROM instructor
```

```
    WHERE dept_name = 'CS')
```

select those instructors having the  
same salary as someone in the  
computer science department

# nested queries, derived relations, and views

## retrieving data using SQL

---

still many other set operations possible:

```
SELECT name
```

```
FROM instructor
```

```
WHERE EXISTS(
```

```
    SELECT name
```

```
    FROM student
```

```
        JOIN advisor ON advisor.student_id = student.id,
```

```
        advisor.instructor_id = instructor.id
```

```
    WHERE tot_cred < 10)
```

select those instructors that  
are advisors to students with  
less than 10 credits in total

**EXISTS** returns **true** if the subquery is not empty

# nested queries, derived relations, and views

## retrieving data using SQL

---

many set operations are possible, but most people will spend years working with databases before they find one “in the wild”

- often (always?) you can write a query in a different way that does not involve set operations
- operations such as **IN(...)**, **EXISTS(...)** are important, as they can be non-trivial to replace
- operations such as **>SOME(...)**, **<>ALL(...)** are seldom used; often easier to solve in alternative ways

you need to be able to write queries with **IN/EXISTS**, but only need to be able to explain what a query with **>some(...)**, **<>all(...)** does

# nested queries, derived relations, and views

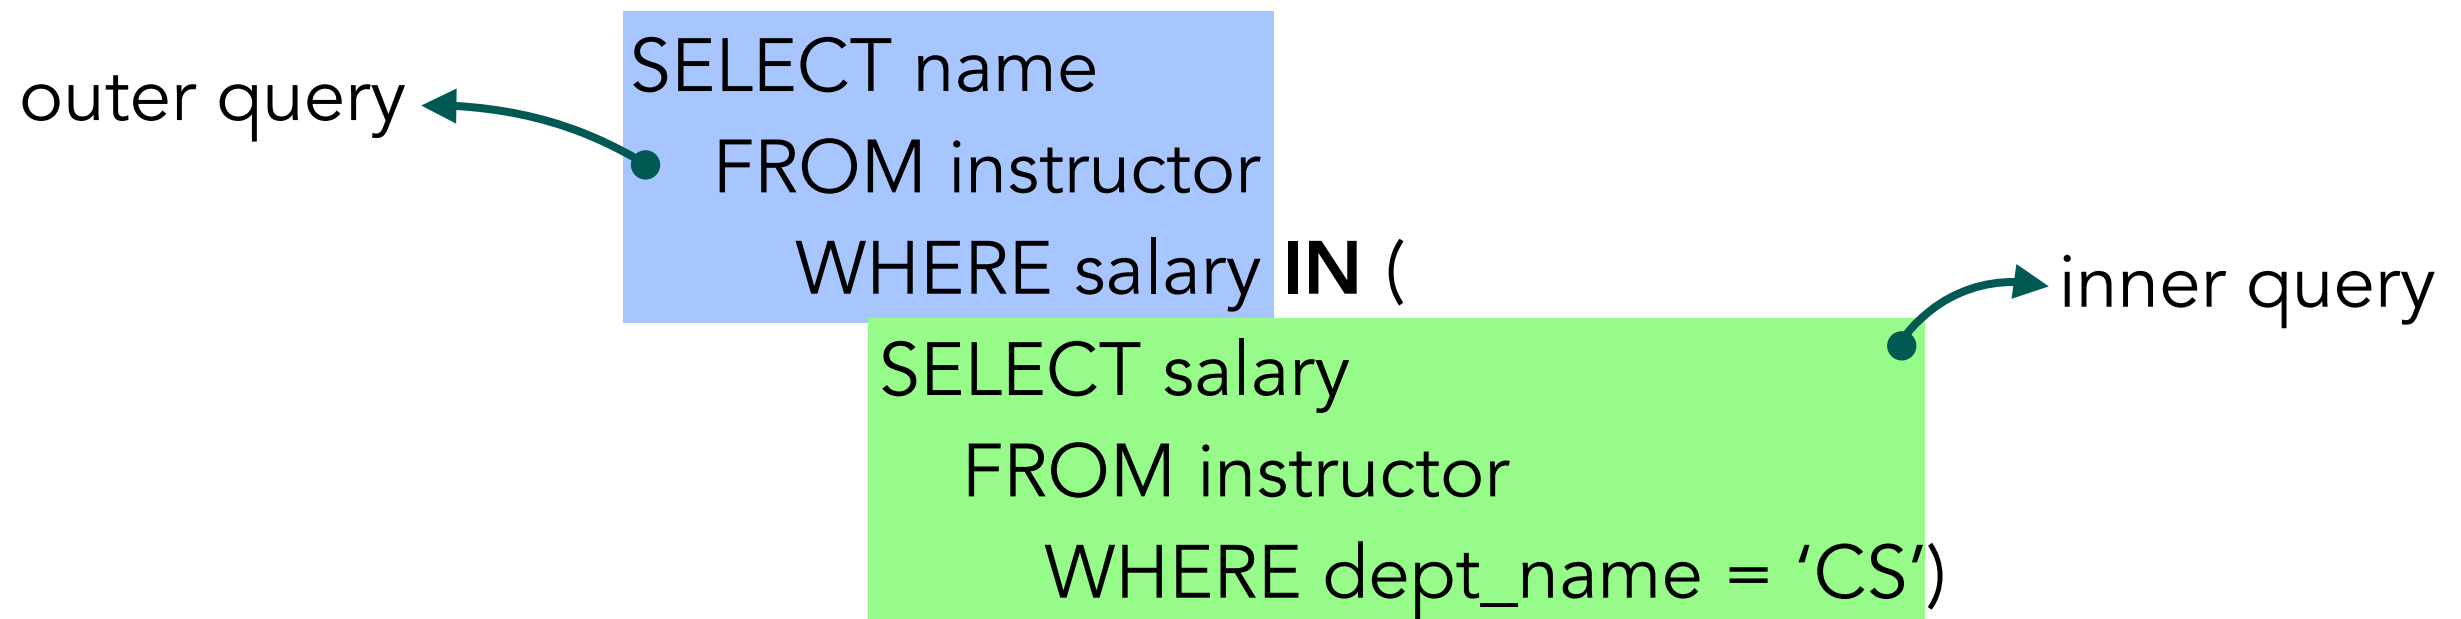
retrieving data using SQL



all set operations so far are examples of nested queries

what is a nested query?

a *nested query*, or *subquery*, or *inner query*, is a query embedded in **parentheses** in the **WHERE** clause of **another query**



# nested queries, derived relations, and views



## retrieving data using SQL

---

nesting queries with SQL

possible to have a query in a query

the inner query only allows a limited syntax

but SELECT ... FROM ... WHERE is generally fine

typically used for set membership/comparison/cardinality

```
SELECT DISTINCT course_id
FROM section
WHERE year = 2014 and
      course_id IN (SELECT course_id
                    FROM section
                    WHERE year = 2015)
```

# nested queries, derived relations, and views

## retrieving data using SQL

---

nesting queries allow us to solve some very hard problems

### problem revisited

what are the names of the students who scored more than the average total credits among the students in the computer science department?

```
SELECT name, tot_cred
FROM student
WHERE tot_cred > (SELECT AVG(tot_cred)
                  FROM student AS inner
                  WHERE inner.dept_name = 'CS')
```

# nested queries, derived relations, and views

## retrieving data using SQL

nested queries have their own terminology

```
SELECT name
FROM instructor
WHERE EXISTS (
    SELECT name
    FROM student
    JOIN advisor ON advisor.student_id = student.id,
                  advisor.instructor_id = instructor.id
    WHERE tot_cred < 10)
```

**correlation** name/variable  
the subquery uses values from the outer query

called a **correlated** subquery  
more on subqueries in next lessons

# nested queries, derived relations, and views



## types of nested queries

---

not all nested queries are created equal!

- (simple) nested subqueries: the inner query can be fully evaluated, and passes results on to outer query
  - ↳ often heavily optimised, and thus very fast  
can be as fast, or faster, than joins
- correlated nested subqueries: inner query uses names/variables from outer query and is therefore reliant on the (results of the) outer query
  - ↳ not a lot of ways to optimise  
typically slow, *sometimes* easiest for complex problems



# nested queries, derived relations, and views



## types of nested queries

---

how are nested subqueries executed?

- (simple) nested subqueries
  - ↳ the innermost query is executed first (and is fully executed)  
the result is passed on to the outer query
- correlated nested subqueries
  - ↳ innermost query is executed *once for each row* in outer query  
needed because inner query references row of outer query

# nested queries, derived relations, and views



## types of nested queries

not all nested queries are created equal!

- single row subquery `SELECT AVG(tot_cred) ...`  
returns a single row, consisting of one column, to the outer query
- multiple row subquery `SELECT name FROM student ...`  
returns multiple rows to the outer query
- multiple column subquery  
returns multiple columns to the outer query

# nested queries, derived relations, and views

single row subquery: use with arithmetic operators

---

single row subquery

↳ the inner query returns only one result to the outer query

```
SELECT name, tot_cred
FROM student
WHERE tot_cred > (SELECT AVG(tot_cred)
                  FROM student AS inner
                  WHERE inner.dept_name = 'CS')
```

can be used with the arithmetic operators =, >, <, >=, <=, < >

# nested queries, derived relations, and views

multiple row subquery: use with multiple value operators

---

multiple row subquery

↳ the inner query returns one or more results to the outer query

```
SELECT DISTINCT course_id
FROM section
WHERE year = 2014 and
      course_id IN (SELECT course_id
                    FROM section
                    WHERE year = 2015)
```

can only be used with multiple value operators,  
such as IN, ANY, ALL, SOME, EXISTS

# nested queries, derived relations, and views

## combinations with arithmetic operators

---

**remember:** can combine ALL/ANY/SOME with arithmetic operators

$$5 > \text{ALL} \left( \begin{array}{|c|} \hline 2 \\ \hline 5 \\ \hline 6 \\ \hline \end{array} \right)$$

*false*, because  $5 \not> 6$

$$5 < \text{ALL} \left( \begin{array}{|c|} \hline 9 \\ \hline 7 \\ \hline 6 \\ \hline \end{array} \right)$$

*true*, because  $5 < 6, 5 < 7, 5 < 9$

$$5 = \text{ALL} \left( \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array} \right)$$

*false*, because  $5 \neq 0$

$$5 <> \text{ALL} \left( \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array} \right)$$


*false*, because  $5 \not<> 5$

# nested queries, derived relations, and views

your turn

try some queries yourself!


- who are the instructors earning more than the average?
- who are the instructors in CS earning more than average for the instructors in PS?
- who are the instructors in CS earning more than any/some/at least one of the instructors in PS?
- who are the instructors working in both PS and CS? (use EXISTS)



# nested queries, derived relations, and views

your turn

who are the instructors in CS not teaching any course?



# nested queries, derived relations, and views

multiple column subquery: the rare one

---

multiple column subquery

↳ the inner query returns one or more tuples to the outer query, each of which consists of more than one column

```
SELECT *  
  FROM section  
  WHERE  
    (building, room_number) IN  
    (SELECT building, room_number  
      FROM classroom  
      WHERE capacity > 100)
```

specify the names of the columns using the ({name, ...}) IN syntax



# nested queries, derived relations, and views

multiple column subquery: the rare one

---

how about retrieving the actual titles of these courses?

```
SELECT title
FROM section
JOIN course USING(course_ID)
WHERE
    (building, room_number) IN
    (SELECT building, room_number
     FROM classroom
     WHERE capacity > 100)
```

multiple column subqueries are not universally supported;  
safer (and easier) to stick to one column

# nested queries, derived relations, and views

subqueries in other parts of queries

---

also possible to have subqueries in HAVING

recall that HAVING is used to select results from a GROUP BY

## problem

which departments have an average salary higher than the average salary of the department of computer science?

```
SELECT dept_name, AVG(salary) AS average
FROM instructor
GROUP BY dept_name
HAVING average > (SELECT AVG(salary)
                  FROM instructor
                  WHERE dept_name = 'Comp. Sci.')
```

# nested queries, derived relations, and views

## derived relations

---

also possible to have subqueries in FROM

↳ called a *derived relation*

these create temporary tables,  
and are especially useful when dealing with groupings!

### problem

give an overview of the departments, as well as the average salary in each department

# nested queries, derived relations, and views

## derived relations

---

also possible to have subqueries in FROM

↳ called a *derived relation*

these create temporary tables,  
and are especially useful when dealing with groupings!

### problem

give an overview of the departments, as well as the average salary in each department

# nested queries, derived relations, and views

## derived relations

---

also possible to have subqueries in FROM

### problem

give an overview of the departments, as well as the average salary in each department

```
SELECT dept_name, avg_salary
FROM (SELECT dept_name, AVG(salary) AS avg_salary
      FROM instructor
      GROUP BY dept_name)
```

can be convenient, as we did not have to use the HAVING clause

# nested queries, derived relations, and views

## derived relations

---

also possible to have subqueries in FROM

### problem

give an overview of the departments and their average salary for those departments with an average salary greater than 80000.

```
SELECT dept_name, avg_salary
FROM (SELECT dept_name, AVG(salary) AS avg_salary
      FROM instructor
      GROUP BY dept_name)
WHERE avg_salary > 80000
```

can be convenient, as we did not have to use the HAVING clause

# nested queries, derived relations, and views

## derived relations

---

but it isn't always needed ...

### problem

give an overview of the departments and their average salary for those departments with an average salary greater than 80000.

```
SELECT dept_name, avg(salary) AS avg_salary  
FROM instructor  
GROUP BY dept_name  
HAVING avg_salary > 80000
```

often more ways than one to achieve your goal in SQL!

# nested queries, derived relations, and views

retrieving data using SQL

---

don't forget:

```
SELECT {attribute [AS new_attribute_name]}  
      FROM {table [AS new_table_name]}  
          [{JOIN table ON {attribute = attribute}}]  
          [WHERE {condition}]  
          [GROUP BY {attribute}]  
          [HAVING {condition}]  
          [ORDER BY {attribute}]
```

where {argument} denotes you need to have at least one, and  
where [argument] denotes a part that is optional and can be omitted



# nested queries, derived relations, and views



why do we need views?

---

remember that security is a reason for using DB

how do we deal with limited access?

access to all instructors, **but not** to their salary

remember that efficiency is a reason for using DB

what if someone always looks at average department salaries?

what if someone always joins multiple tables?

*duplication of table is a bad idea!*

views are stored queries that can be used like any other table

*like a variable in programming!*

# nested queries, derived relations, and views



## creating views

a view is any relation visible to a user that is not a conceptual model

CREATE VIEW view\_name AS sql\_expression

name you choose ←

→ any valid SQL expression

once a view is created, *view\_name* can be used as any other table

crucially, a view does **not** create a new table

defining a view saves the SQL expression,  
and substitutes that expression as needed in other queries

# nested queries, derived relations, and views

## example

---

create a view that excludes salary:

```
CREATE VIEW faculty AS  
  SELECT id, name, dept_name  
  FROM instructor
```

```
SELECT * FROM faculty
```

*result:*

<u>id</u>	name	dept_name
45039	John	Comp. Sci.
30594	Selma	Comp. Sci.
30492	Paul	Elec. Eng.
20996	Tisan	Elec. Eng.


# nested queries, derived relations, and views

your turn

---

try creating some views yourself!

- create a view of all instructors on the Biology department
- create a view of department salary totals



# nested queries, derived relations, and views

## views defined on views

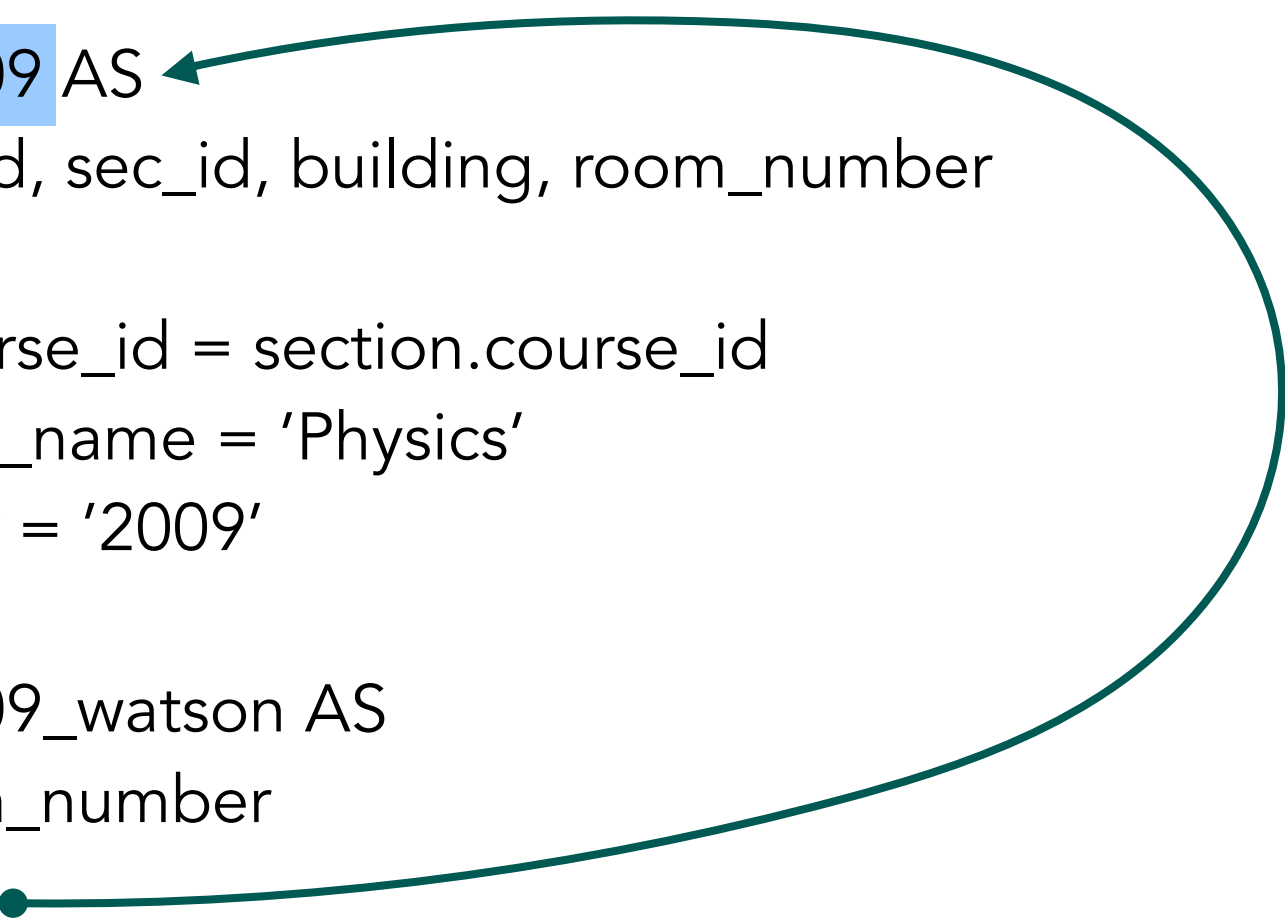
---

views defined on views

*remember:* once a view is created, it acts like a table

```
CREATE VIEW physics_2009 AS  
  SELECT course.course_id, sec_id, building, room_number  
  FROM course, section  
  WHERE course.course_id = section.course_id  
    AND course.dept_name = 'Physics'  
    AND section.year = '2009'
```

```
CREATE VIEW physics_2009_watson AS  
  SELECT course_id, room_number  
  FROM physics_2009  
  WHERE building = 'Watson'
```



# nested queries, derived relations, and views

## views defined on views: terminology

---

views defined on views

*some terminology:*

a view  $v$  **depends directly** on  $w$  when  $w$  is used in the definition of the view  $v$

```
CREATE VIEW w ...
```

```
CREATE VIEW v ...  
  SELECT ...  
  FROM w
```

a view  $v$  **depends** on  $w$  if  $v$  depends directly on  $w$ , or some view used (indirectly) by  $v$  depends directly on  $w$

a view  $v$  is **recursive** if it depends on itself, i.e. view  $v$

# nested queries, derived relations, and views

views as “variables”

---

using *with* to define a temporary view

when creating a difficult query, we can temporarily save and use a result as a view, accessible only by that query

```
WITH max_budget (value) AS
  (SELECT max(budget)
   FROM department)
SELECT budget
  FROM department, max_budget
 WHERE department.budget = max_budget.value
```

# nested queries, derived relations, and views

## updating views

---

since views are virtual tables, we can update them

```
CREATE VIEW faculty AS  
  SELECT id, name, dept_name  
  FROM instructor
```

```
INSERT INTO faculty VALUES ('30765', 'Green', 'Music')
```

corresponds to:

```
INSERT INTO instructor VALUES ('30765', 'Green', 'Music', null)
```



# nested queries, derived relations, and views

## updating views

---

since views are virtual tables, we can update them **to some extent**

```
CREATE VIEW instructor_info AS  
  SELECT ID, name, building  
  FROM instructor  
  JOIN department USING(dept_name)
```

```
INSERT INTO instructor_info VALUES ('69987', 'White', 'Taylor')
```

some unresolved ambiguity:

- what if there are no departments located in building Taylor?
- what if there are multiple departments in Taylor?

# nested queries, derived relations, and views

## updating views

---



updating generally only allowed on *updatable* views:

1. The FROM clause has only one database relation.
2. The SELECT clause contains only attribute names of the relation, and does not have any expressions, aggregates, or DISTINCT specification.
3. Any attribute not listed in the SELECT clause can be set to *null*.
4. The query does not have a GROUP BY or HAVING clause.